

# Yijie Qiu

helen.yijieqiu@gmail.com | <https://www.linkedin.com/in/yijieqiu> | 415.823.1986

## EDUCATION

**CARNEGIE MELLON UNIVERSITY**  
MS INFORMATION SYSTEMS MGMT  
December 2015 | Pittsburgh, PA

**BS CHEMICAL ENGINEERING**  
May 2014 | Pittsburgh, PA  
Additional major in Economics and Statistics  
University Honors

## COURSEWORK

### GRADUATE

10605 Machine Learning Large Data Sets  
11642 Search Engine  
15619 Cloud Computing  
95702 Distributed Systems  
95877 Python for Data Analysis  
95736 Advanced Relational Database

### UNDERGRADUATE

15121 Data Structures in Java  
36402 Advanced Data Analysis  
36315 Statistical Graphics and Visualization  
06462 Optimization Modeling & Algorithms

## SKILL

### TECHNICAL

#### PROGRAMMING LANGUAGES

Java  
Python  
Ruby  
Shell  
R  
MATLAB

#### TOOLS

AWS Services  
Oracle Database & PL/SQL  
Hive  
MapReduce

### LANGUAGE

Native: Mandarin Chinese  
Fluent: English • Cantonese

## EXPERIENCE

**AMAZON | SOFTWARE ENGINEER, AMAZON SAGEMAKER HOSTING**  
March 2017 - Present | Seattle, WA

*Amazon SageMaker is a fully-managed machine learning platform that enables developers and data scientists to build, train, and deploy ML models at scale. The hosting service provides managed HTTPS endpoints, which let users go from ML model Docker containers to production-ready, scalable online inferences in just a few clicks.*

- Contributed to initial launch of Amazon SageMaker at AWS re:Invent 2017.
- Designed and implemented APIs and workflows for bouncing the underlying infrastructure of SageMaker endpoints without interruption to inferences.
- Architected and implemented system for automated detection of SageMaker endpoints running outdated software, and applying necessary patches with zero downtime. The system has orchestrated software and security updates for more than 15K endpoints to-date.
- Implemented service priming mechanism for SageMaker online inference routing fleet, to minimize cold start impact on server-side latency. For requests with 1KB payload, the P99 overhead latency immediately following service deployment dropped from >1.6 second to 80 ms with priming.
- Built service canaries for all SageMaker Hosting APIs with extensible design, allowing for continuous monitoring of service health and requiring minimal code changes to add coverage for new features.
- Improved service scalability and lowered endpoint operation latency by incorporating event-driven architectures into asynchronous workflows.

**AMAZON | SOFTWARE ENGINEER, MOBILE DIAGNOSTICS**  
February 2016 - March 2017 | Seattle, WA

*Part of Amazon device middleware services, the Mobile Diagnostics team enables Amazon devices to record key app performance, usage, and crash information, and securely transmit them to backend services for troubleshooting and analytics purposes.*

- Developed an end-to-end encrypted, terabyte scale data ingestion pipeline. The system transforms app and devices usage metrics collected on individual hardware (>280 million entries daily as of mid-2017) into structured, aggregated BI tables capable of providing consumer insights and assisting business decision making.
- Established load testing utilities for back-end services, and instrumented load tests and server, load balancer tuning in preparation of holiday peaks.
- Maintained server-side and client-side metrics and logging components used by Amazon apps across various platforms (Fire OS, Android, iOS).

**AMAZON | SOFTWARE DEVELOPMENT INTERN, KINDLE XRAY**  
May 2015 - Aug 2015 | Seattle, WA

*Kindle X-ray is a content enhancement feature that notes the appearances of key characters, places, and phrases throughout the book, and provides detailed descriptions, all with a single tap.*

- Performed analysis using MapReduce to identify areas of improvement in text labeling pipelines.
- Improved quality of pipeline outputs by implementing and integrating with an AWS RDS database containing historic labelling and correction data.